

dip

February 10, 2025

```
[33]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import adjusted_rand_score

from sklearn.cluster import KMeans

from sklearn.decomposition import PCA

# Load dataset
#file_path = "/mnt/data/mcdonalds.csv"
mcdonalds = pd.read_csv("mcdonalds.csv")

# Inspect dataset
print("Column Names:", mcdonalds.columns.tolist())
print("Dataset Shape:", mcdonalds.shape)
print("First Three Rows:\n", mcdonalds.head(3))

# Convert segmentation variables from Yes/No to binary numeric
MD_x = mcdonalds.iloc[:, :11].applymap(lambda x: 1 if x == "Yes" else 0)
print("\nAverage values of transformed variables:\n", MD_x.mean().round(2))

# Perform Principal Component Analysis (PCA)
pca = PCA()
MD_pca = pca.fit_transform(MD_x)

# Display explained variance
explained_variance = np.round(pca.explained_variance_ratio_, 4)
cumulative_variance = np.cumsum(explained_variance)
print("\nExplained Variance by Principal Components:\n", explained_variance)
print("Cumulative Variance:\n", cumulative_variance)

# Display principal component loadings
loadings = pd.DataFrame(pca.components_.T, index=MD_x.columns,
    columns=[f'PC{i+1}' for i in range(MD_x.shape[1])])
print("\nPCA Loadings:\n", loadings.round(2))
```

```

# Plot Perceptual Map (First Two Principal Components)
plt.figure(figsize=(8, 6))
plt.scatter(MD_pca[:, 0], MD_pca[:, 1], color='grey', alpha=0.5)
for i, txt in enumerate(MD_x.columns):
    plt.arrow(0, 0, loadings.iloc[i, 0]*2, loadings.iloc[i, 1]*2, color='red',
    head_width=0.05)
    plt.text(loadings.iloc[i, 0]*2.2, loadings.iloc[i, 1]*2.2, txt,
    color='black')
plt.axhline(0, color='black', linewidth=0.5)
plt.axvline(0, color='black', linewidth=0.5)
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.title("Perceptual Map of McDonald's Attributes")
plt.grid()
plt.show()

```

Column Names: ['yummy', 'convenient', 'spicy', 'fattening', 'greasy', 'fast', 'cheap', 'tasty', 'expensive', 'healthy', 'disgusting', 'Like', 'Age', 'VisitFrequency', 'Gender']

Dataset Shape: (1453, 15)

First Three Rows:

	yummy	convenient	spicy	fattening	greasy	fast	cheap	tasty	expensive	healthy	\
0	No	Yes	No	Yes	No	Yes	Yes	No	Yes	No	
1	Yes	Yes	No	Yes	Yes	Yes	Yes	Yes	Yes	No	
2	No	Yes	Yes	Yes	Yes	Yes	No	Yes	Yes	Yes	

	disgusting	Like	Age	VisitFrequency	Gender
0	No	-3	61	Every three months	Female
1	No	+2	51	Every three months	Female
2	No	+1	62	Every three months	Female

Average values of transformed variables:

yummy	0.55
convenient	0.91
spicy	0.09
fattening	0.87
greasy	0.53
fast	0.90
cheap	0.60
tasty	0.64
expensive	0.36
healthy	0.20
disgusting	0.24

dtype: float64

Explained Variance by Principal Components:

[0.2994 0.1928 0.133 0.0831 0.0595 0.0503 0.0438 0.0395 0.0368 0.0324

0.0293]

Cumulative Variance:


[0.2994 0.4922 0.6252 0.7083 0.7678 0.8181 0.8619 0.9014 0.9382 0.9706
0.9999]

PCA Loadings:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9	PC10	PC11
yummy	-0.48	0.36	-0.30	0.06	-0.31	0.17	-0.28	0.01	0.57	-0.11	0.05
convenient	-0.16	0.02	-0.06	-0.14	0.28	-0.35	-0.06	-0.11	-0.02	-0.67	-0.54
spicy	-0.01	0.02	-0.04	0.20	0.07	-0.36	0.71	0.38	0.40	-0.08	0.14
fattening	0.12	-0.03	-0.32	-0.35	-0.07	-0.41	-0.39	0.59	-0.16	-0.01	0.25
greasy	0.30	-0.06	-0.80	0.25	0.36	0.21	0.04	-0.14	-0.00	0.01	0.00
fast	-0.11	-0.09	-0.06	-0.10	0.11	-0.59	-0.09	-0.63	0.17	0.24	0.34
cheap	-0.34	-0.61	-0.15	0.12	-0.13	-0.10	-0.04	0.14	0.08	0.43	-0.49
tasty	-0.47	0.31	-0.29	-0.00	-0.21	-0.08	0.36	-0.07	-0.64	0.08	0.02
expensive	0.33	0.60	0.02	0.07	-0.00	-0.26	-0.07	0.03	0.07	0.45	-0.49
healthy	-0.21	0.08	0.19	0.76	0.29	-0.18	-0.35	0.18	-0.19	-0.04	0.16
disgusting	0.37	-0.14	-0.09	0.37	-0.73	-0.21	-0.03	-0.17	-0.07	-0.29	-0.04

/var/folders/ns/_gksq54x6hq35t894m1570280000gn/T/ipykernel_5408/2757169144.py:20
: FutureWarning: DataFrame.applymap has been deprecated. Use DataFrame.map
instead.

```
MD_x = mcdonalds.iloc[:, :11].applymap(lambda x: 1 if x == "Yes" else 0)
```



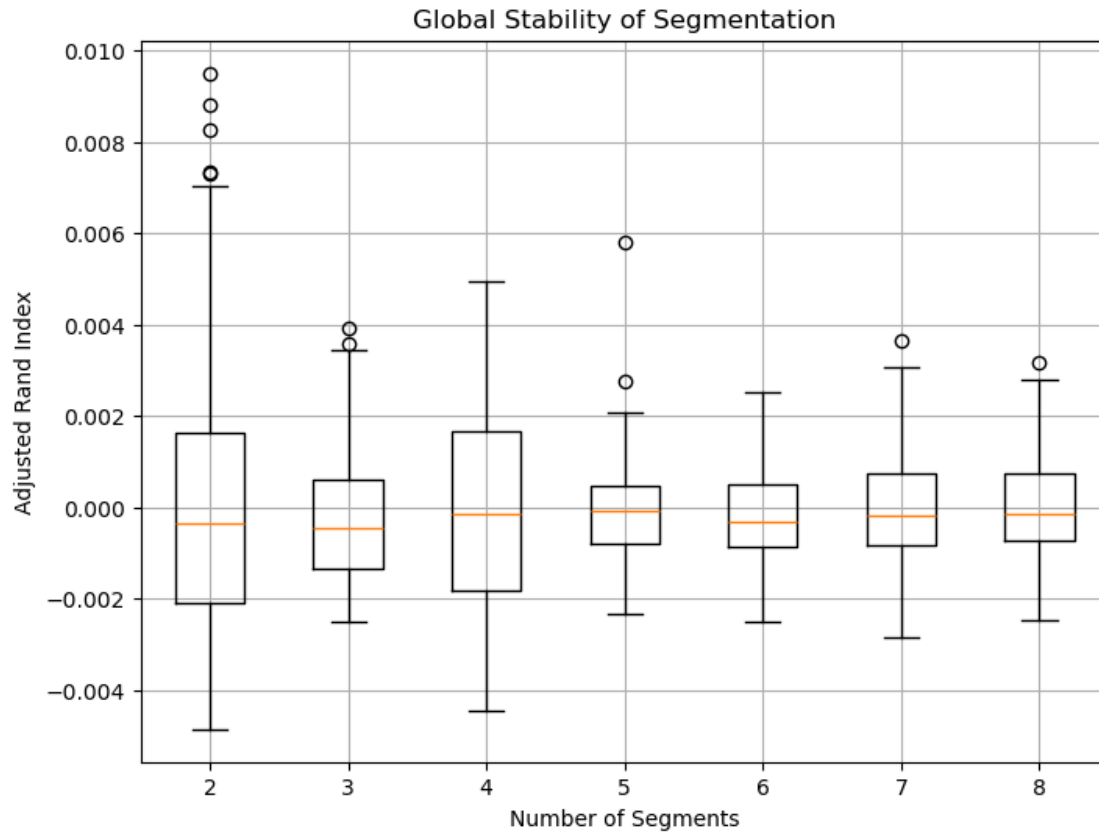
```

[26]: # Stability analysis using K-Means clustering
num_segments = range(2, 9)
n_boot = 100
stability_scores = []

for k in num_segments:
    rand_indices = []
    for _ in range(n_boot):
        sample_idx = np.random.choice(len(MD_x), len(MD_x), replace=True)
        sample_data = MD_x.iloc[sample_idx]
        kmeans = KMeans(n_clusters=k, n_init=10, random_state=None).
        ↪fit(sample_data)
        ref_kmeans = KMeans(n_clusters=k, n_init=10, random_state=None).
        ↪fit(MD_x)
        rand_indices.append(adjusted_rand_score(ref_kmeans.labels_, kmeans.
        ↪labels_))
    stability_scores.append(rand_indices)

# Boxplot of stability scores
plt.figure(figsize=(8, 6))
plt.boxplot(stability_scores, labels=num_segments)
plt.xlabel("Number of Segments")
plt.ylabel("Adjusted Rand Index")
plt.title("Global Stability of Segmentation")
plt.grid()
plt.show()

```

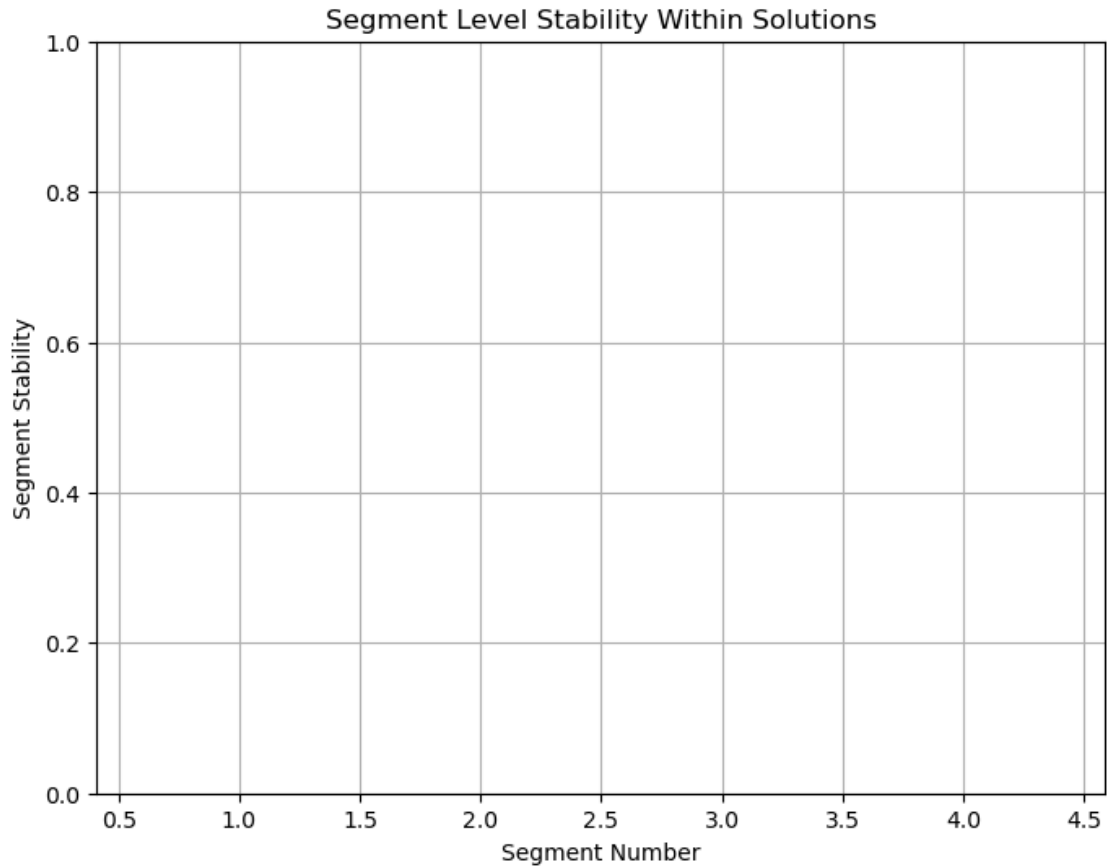


```
[37]: # Selecting the best segmentation solution
best_k = 4
final_kmeans = KMeans(n_clusters=best_k, n_init=10, random_state=42).fit(MD_x)
mcdonalds["Segment"] = final_kmeans.labels_

# Segment Level Stability Within Solutions (SLSW)
segment_stability = []
for segment in range(best_k):
    segment_indices = (final_kmeans.labels_ == segment)
    intra_stability = np.mean([adjusted_rand_score(final_kmeans.
↪ labels_[segment_indices], KMeans(n_clusters=best_k, n_init=10,
↪ random_state=None).fit(MD_x.iloc[segment_indices]).labels_) for _ in
↪ range(n_boot)])
    segment_stability.append(intra_stability)

# Plot segment stability
plt.figure(figsize=(8, 6))
plt.bar(range(1, best_k+1), segment_stability)
plt.xlabel("Segment Number")
plt.ylabel("Segment Stability")
```

```
plt.title("Segment Level Stability Within Solutions")
plt.ylim(0, 1)
plt.grid()
plt.show()
```



```
[41]: # Convert segmentation variables from Yes/No to binary numeric
MD_x = mcdonalds.iloc[:, :11].applymap(lambda x: 1 if x == "Yes" else 0)

# Fit Gaussian Mixture Model (GMM) for 2 to 8 segments
num_segments = range(2, 9)
gmm_models = {}
log_likelihoods = []
aic_values = []
bic_values = []

for k in num_segments:
    gmm = GaussianMixture(n_components=k, n_init=10, random_state=1234)
    gmm.fit(MD_x)
    gmm_models[k] = gmm
```

```

log_likelihoods.append(gmm.lower_bound_ * MD_x.shape[0])
aic_values.append(gmm.aic(MD_x))
bic_values.append(gmm.bic(MD_x))

# Plot Information Criteria
plt.figure(figsize=(8, 6))
plt.plot(num_segments, aic_values, label='AIC', marker='o')
plt.plot(num_segments, bic_values, label='BIC', marker='s')
plt.xlabel("Number of Segments")
plt.ylabel("Information Criteria (AIC, BIC)")
plt.title("Model Selection using Information Criteria")
plt.legend()
plt.grid()
plt.show()

# Select best number of segments (based on elbow method, e.g., 4)
best_k = 4
best_gmm = gmm_models[best_k]
mcdonalds["GMM_Segment"] = best_gmm.predict(MD_x)

# Compare with K-Means clustering
kmeans = KMeans(n_clusters=best_k, n_init=10, random_state=42).fit(MD_x)
mcdonalds["KMeans_Segment"] = kmeans.labels_

# Cross-tabulation of segments
cross_tab = pd.crosstab(mcdonalds["KMeans_Segment"], mcdonalds["GMM_Segment"])
print("\nCross-tabulation of K-Means and GMM Segments:\n", cross_tab)

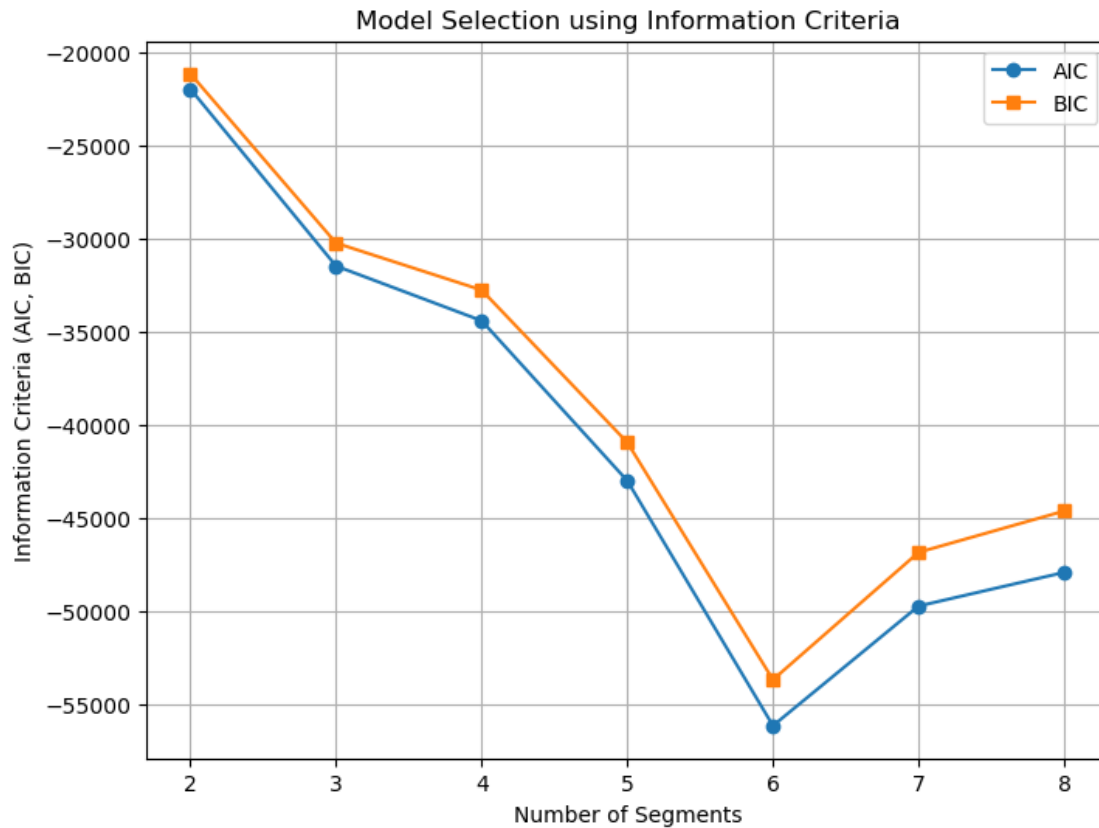
# Initializing GMM using K-Means clusters
gmm_kmeans_init = GaussianMixture(n_components=best_k, n_init=10,
    random_state=1234, init_params='kmeans')
gmm_kmeans_init.fit(MD_x)
mcdonalds["GMM_KMeansInit_Segment"] = gmm_kmeans_init.predict(MD_x)

# Compare log-likelihood values
print("\nLog-Likelihood Values:")
print("Random Initialization:", best_gmm.lower_bound_ * MD_x.shape[0])
print("K-Means Initialization:", gmm_kmeans_init.lower_bound_ * MD_x.shape[0])

```

/var/folders/ns/_gksq54x6hq35t894m1570280000gn/T/ipykernel_5408/1934968170.py:2:
FutureWarning: DataFrame.applymap has been deprecated. Use DataFrame.map
instead.

```
MD_x = mcdonalds.iloc[:, :11].applymap(lambda x: 1 if x == "Yes" else 0)
```



Cross-tabulation of K-Means and GMM Segments:

GMM_Segment	0	1	2	3
KMeans_Segment				
0	4	213	0	23
1	120	0	403	26
2	26	52	227	4
3	28	3	0	324

Log-Likelihood Values:

Random Initialization: 17513.53905912311

K-Means Initialization: 17513.53905912311

```
[47]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import statsmodels.api as sm
from sklearn.mixture import GaussianMixture
```



```

# Mapping for 'Like' column to numerical values
like_mapping = {
    "I love it!+5": 5, "+4": 4, "+3": 3, "+2": 2, "+1": 1, "0": 0,
    "-1": -1, "-2": -2, "-3": -3, "-4": -4, "I hate it!-5": -5
}

# Convert categorical 'Like' variable to numerical scale (-5 to +5)
mcdonalds['Like.n'] = mcdonalds['Like'].map(like_mapping)

# Drop rows where Like.n could not be mapped
mcdonalds = mcdonalds.dropna(subset=['Like.n'])

# Define dependent variable (y) and independent variables (X)
X = mcdonalds.iloc[:, :11].applymap(lambda x: 1 if x == "Yes" else 0)
y = mcdonalds['Like.n']

# Fit a mixture of regression models with 2 components
num_components = 2
gmm = GaussianMixture(n_components=num_components, n_init=10, random_state=1234)
clusters = gmm.fit_predict(X)
mcdonalds['GMM_Cluster'] = clusters

# Fit separate regression models for each segment
regression_results = {}
for segment in range(num_components):
    segment_data = mcdonalds[mcdonalds['GMM_Cluster'] == segment]
    X_segment = sm.add_constant(segment_data.iloc[:, :11].applymap(lambda x: 1
↪if x == "Yes" else 0))
    y_segment = segment_data['Like.n']
    model = sm.OLS(y_segment, X_segment).fit()
    regression_results[segment] = model
    print(f"\nSegment {segment + 1} Regression Results:")
    print(model.summary())

# Extract regression coefficients for plotting
coefficients = pd.DataFrame({f'Segment {i+1}': model.params for i, model in
↪regression_results.items()})
coefficients = coefficients.drop("const")

# Plot regression coefficients
plt.figure(figsize=(10, 6))
coefficients.plot(kind='bar', figsize=(12, 6), alpha=0.75)
plt.axhline(y=0, color='black', linewidth=0.75, linestyle='--')
plt.xlabel("Independent Variables")
plt.ylabel("Regression Coefficients")
plt.title("Comparison of Regression Coefficients Across Segments")
plt.xticks(rotation=45)

```

```
plt.legend()
plt.grid()
plt.show()
```

```
/var/folders/ns/_gksq54x6hq35t894m1570280000gn/T/ipykernel_5408/1622017903.py:22
: FutureWarning: DataFrame.applymap has been deprecated. Use DataFrame.map
instead.
```

```
X = mcdonalds.iloc[:, :11].applymap(lambda x: 1 if x == "Yes" else 0)
/var/folders/ns/_gksq54x6hq35t894m1570280000gn/T/ipykernel_5408/1622017903.py:35
: FutureWarning: DataFrame.applymap has been deprecated. Use DataFrame.map
instead.
```

```
X_segment = sm.add_constant(segment_data.iloc[:, :11].applymap(lambda x: 1 if
x == "Yes" else 0))
```

```
/opt/anaconda3/lib/python3.12/site-
```

```
packages/statsmodels/regression/linear_model.py:1967: RuntimeWarning: divide by
zero encountered in scalar divide
```

```
    return np.sqrt(eigvals[0]/eigvals[-1])
```

```
/var/folders/ns/_gksq54x6hq35t894m1570280000gn/T/ipykernel_5408/1622017903.py:35
: FutureWarning: DataFrame.applymap has been deprecated. Use DataFrame.map
instead.
```

```
X_segment = sm.add_constant(segment_data.iloc[:, :11].applymap(lambda x: 1 if
x == "Yes" else 0))
```

Segment 1 Regression Results:

OLS Regression Results

```
=====
Dep. Variable:          Like.n      R-squared:          0.405
Model:                  OLS         Adj. R-squared:      0.400
Method:                 Least Squares  F-statistic:        83.00
Date:                  Mon, 10 Feb 2025  Prob (F-statistic):    1.34e-104
Time:                  21:04:37      Log-Likelihood:     -1982.1
No. Observations:      985          AIC:                  3982.
Df Residuals:          976          BIC:                  4026.
Df Model:               8
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
yummy	1.9765	0.160	12.375	0.000	1.663	2.290
convenient	-0.0480	0.125	-0.385	0.701	-0.293	0.197
spicy	-0.5211	0.203	-2.565	0.010	-0.920	-0.122
fattening	-0.5446	0.182	-3.000	0.003	-0.901	-0.188
greasy	-0.3425	0.123	-2.777	0.006	-0.584	-0.100
fast	-0.0480	0.125	-0.385	0.701	-0.293	0.197
cheap	0.0886	0.181	0.489	0.625	-0.267	0.444
tasty	1.4321	0.177	8.077	0.000	1.084	1.780
expensive	0.0891	0.191	0.466	0.641	-0.286	0.464

healthy	0.4674	0.150	3.116	0.002	0.173	0.762
disgusting	0	0	nan	nan	0	0

```
=====
```

Omnibus:	59.066	Durbin-Watson:	2.020
Prob(Omnibus):	0.000	Jarque-Bera (JB):	69.016
Skew:	-0.603	Prob(JB):	1.03e-15
Kurtosis:	3.475	Cond. No.	inf

```
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 0. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Segment 2 Regression Results:

OLS Regression Results

```
=====
```

Dep. Variable:	Like.n	R-squared:	0.588
Model:	OLS	Adj. R-squared:	0.578
Method:	Least Squares	F-statistic:	59.23
Date:	Mon, 10 Feb 2025	Prob (F-statistic):	1.03e-80
Time:	21:04:37	Log-Likelihood:	-1001.6
No. Observations:	468	AIC:	2027.
Df Residuals:	456	BIC:	2077.
Df Model:	11		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-2.4721	0.402	-6.143	0.000	-3.263	-1.681
yummy	3.2251	0.303	10.638	0.000	2.629	3.821
convenient	0.9483	0.226	4.195	0.000	0.504	1.392
spicy	-0.0656	0.343	-0.191	0.849	-0.740	0.609
fattening	0.0229	0.369	0.062	0.951	-0.702	0.748
greasy	-0.0312	0.244	-0.128	0.898	-0.511	0.449
fast	0.4303	0.257	1.673	0.095	-0.075	0.936
cheap	-0.0294	0.282	-0.104	0.917	-0.583	0.525
tasty	1.3174	0.280	4.711	0.000	0.768	1.867
expensive	-0.1235	0.271	-0.455	0.649	-0.657	0.410
healthy	0.3981	0.346	1.151	0.250	-0.282	1.078
disgusting	-1.9007	0.305	-6.233	0.000	-2.500	-1.301

```
=====
```

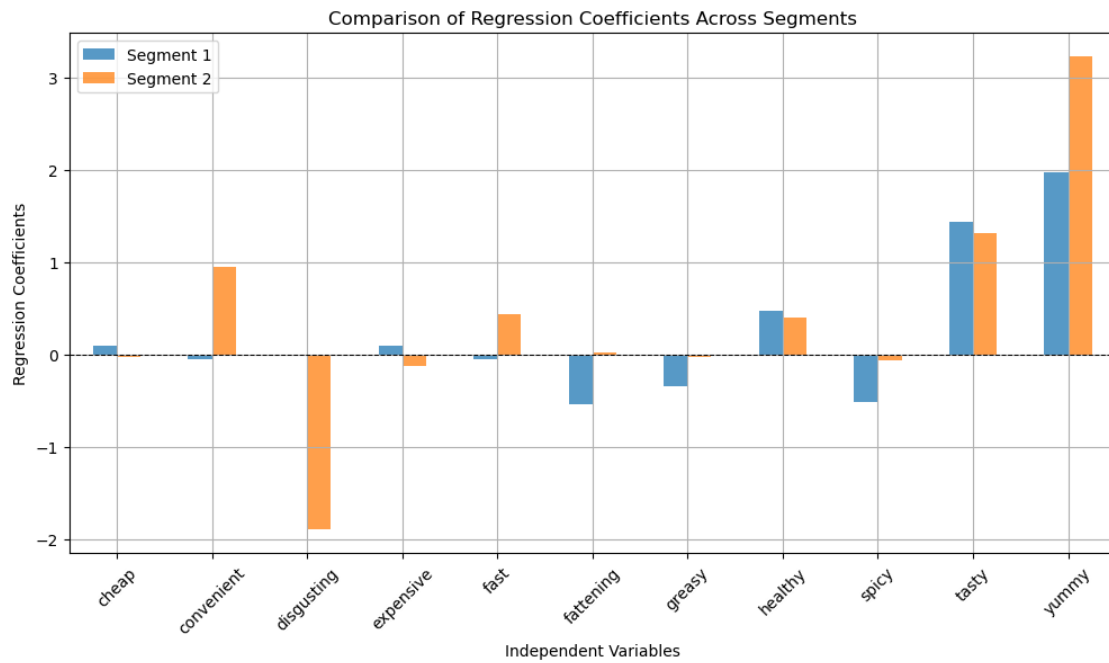
Omnibus:	2.976	Durbin-Watson:	1.909
Prob(Omnibus):	0.226	Jarque-Bera (JB):	2.738
Skew:	0.164	Prob(JB):	0.254
Kurtosis:	3.180	Cond. No.	11.3

```
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

<Figure size 1000x600 with 0 Axes>



[]: